# Design Experience of an SRv6 uSID Data Center

SRv6 uSID DC Landscape

Gyan Mishra, GTS

"IT Technologist & Innovation Specialist"

Associate Fellow –Network Design

April 9<sup>th</sup> 2024

YouTube Video!!



# **Demo time!**

### All Demo's @ YouTube Channel SRv6 uSID DC:

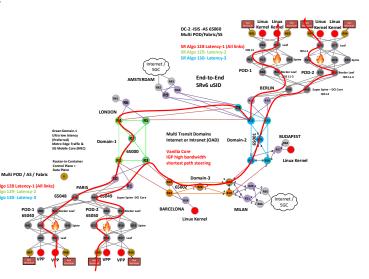
https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

Use-case 2: SRv6 uSID w/ Multi POD Fabrics

Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

- ☐ Host-to-Host multi-pod across the metro
- Policy programmed from Linux Kernel & VPP host
- ☐ Policy programmed from Router-In-Container (XRD) DC.1 Multi POD / AS / Fabric





# Real World Use-cases

**#1 IPv6 Host Based Networking** 

#2 Dual Plane MPLS / IPv6 Core Migration

#3 SRv6 uSID End-To-End



# **#1 IPv6 Host Based Networking**

- Traffic Engineering and Carrier Grade features are not a requirement in the Data Center.
- Operators can use white box switches or disaggregated hardware and software with Vanilla IPv6 Only DC fabric blindly passing the SRv6 uSID packets. Massive bandwidth where Multi Petabits of fiber can be thrown at the DC fabric, with the focus on High Bandwidth packet pushing with Ultra simplified fabric.
- **Steering** is initiated from the Data Center host attachment using IGP shortest path leaving the entire fabric 100% vanilla IPv6.



# #2 Dual Plane MPLS / IPv6 Core Migration

- Traffic Engineering & Carrier Grade features are a requirement ONLY in the Data Center.
- Traffic Engineering capabilities in the Data Center, and the intermediate domains follow IGP shortest path blindly forwarding the SRv6 uSID packets. Massive scale & resiliency with full carrier grade features in the Data Center.
- **Steering** is initiated from the DC host attachment and follows IGP shortest path along the intermediate domains to the egress DC or Domain.



# #3 SRv6 uSID End-To-End

- Traffic Engineering and all the Carrier Grade features are a requirement.
- Full feature richness.
- **Steering** is initiated from the Data Center host attachment or could be any switch within the DC fabric for any and all flow types.



# This has too much on slide so not using - too busy!

# **Demo time!**

All Demo's @ YouTube Channel IPv6 uSID DC:

https://youtube.com/@SRv6\_uSID\_DC

https://github.com/segmentrouting/srv6-labs

- ☐ Host-to-Host multi-pod across the metro
- → Policy programmed from Linux Kernel & VPP host

#### Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)
- SRv6 uSID End-to-End (1c)

#### Use-case 2: SRv6 uSID w/ Multi POD Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (2a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2b)
- SRv6 uSID End-to-End (2c)

#### Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)
- SRv6 uSID End-to-End (3c)

Don't forget to LIKE the video & provide feedback for future Video's on IPv6 DC Arch

# This is & ongoing SRv6 uSID DC R&D work By Contributors:

Gyan Mishra @ Verizon
gyan.s.mishra@verizon.com
Bruce McDougall @ Cisco
brmcdoug@cisco.com



# Demo time!

#### All Demo's @ YouTube Channel IPv6 uSID DC:

https://youtube.com/@SRv6\_uSID\_DC

https://github.com/segmentrouting/srv6-labs

- ☐ Host-to-Host multi-pod across the metro
- Policy programmed from Linux Kernel & VPP host

#### Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)
- SRv6 uSID uSID End-to-End (1c)

#### Use-case 2: SRv6 uSID w/ Multi POD Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (2a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2b)
- SRv6 uSID End-to-End (2c)

#### Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)
- SRv6 uSID End-to-End (3c)

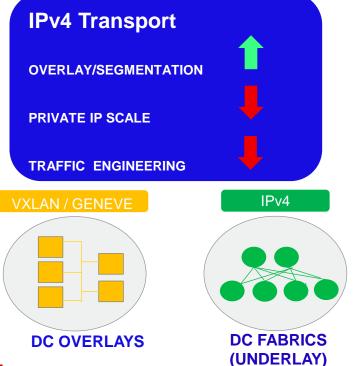
# MPLS WC 2024 SRv6 uSID DC Series

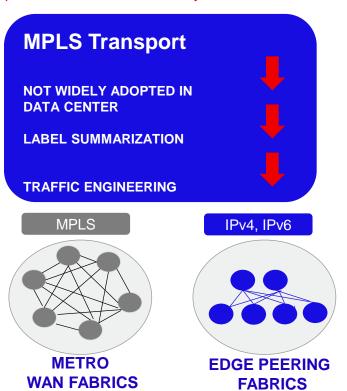
- ⇔ YouTube Video #1 (1 of 9)
- ⇔ YouTube Video #2 (2 of 9)
- ⇔ YouTube Video #3 (3 of 9)
- ⇔ YouTube Video #4 (4 of 9)
- ⇔ YouTube Video #5 (5of 9)
- ⇔ YouTube Video #6 (6 of 9)
- ⇔ YouTube Video #7 (7 of 9)
- ⇔ YouTube Video #8 (8of 9)
- ⇔ YouTube Video #9 (9 of 9)



# SILOED Networking ⇔ Complexity Tax

Each siloed network Domain has its own Hardware, Software, SDN Stack, Operations & Automation Ecosystem





# SRv6 uSID ⇔ Simplicity, Functionality, Ultra Scale

A Common End-to-End Forwarding Architecture Enables Common HW, SW, SDN, Ops ⇔ Massive Scale



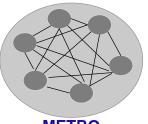
#### SRv6 uSID ⇔ Translates into Ultra Massive Scale



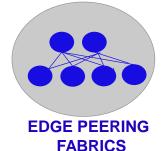


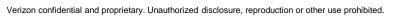
**DC FABRICS** 

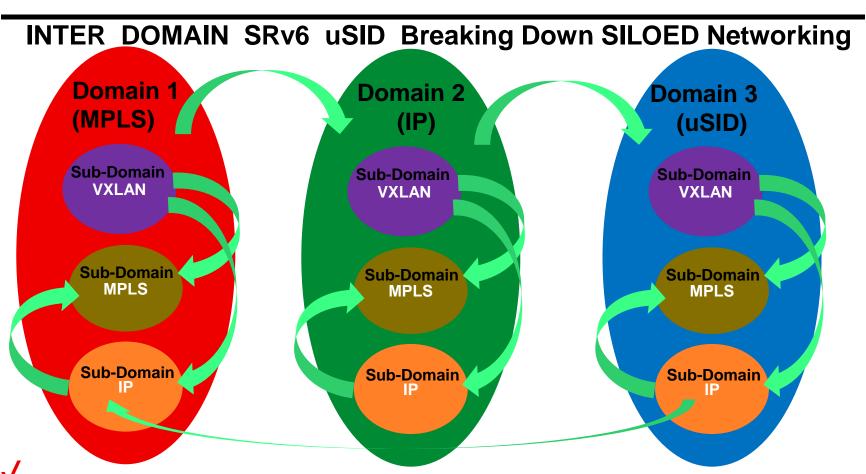
"End-to-End Inter-Domain Routing via SRv6 uSID"

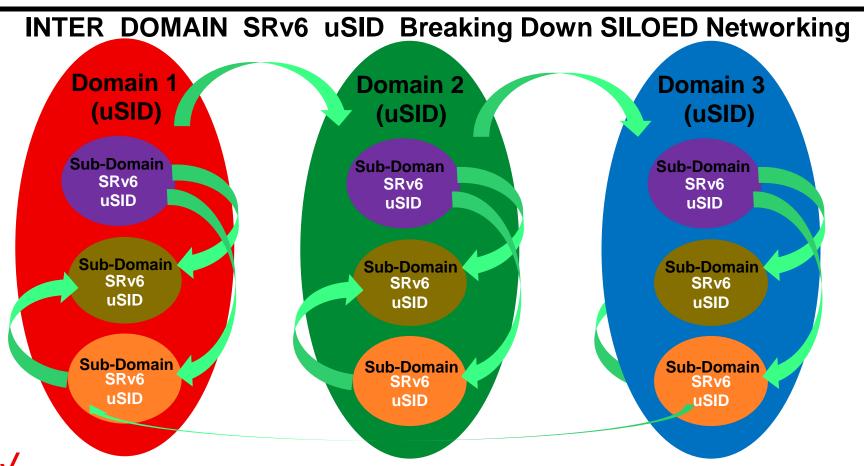


METRO WAN FABRICS

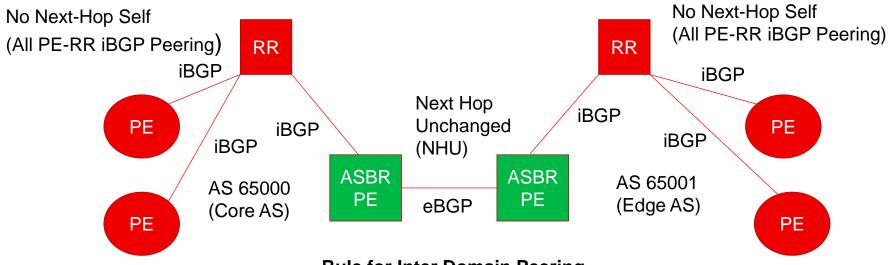








# INTER DOMAIN SRv6 uSID



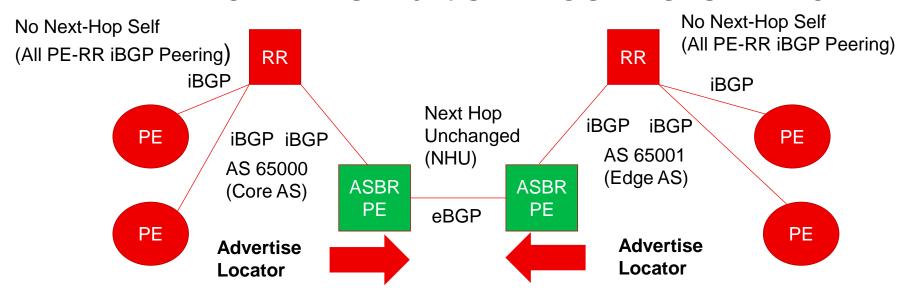
# Rule for Inter Domain Peering

- iBGP -No Next-Hop Self
- eBGP –Next Hop Unchanged

(Requirement to Preserve L2 VPN & L3 VPN Service SID across INTER-AS Boundary)



# INTER DOMAIN SRV6 uSID ROUTING SIMPLICITY



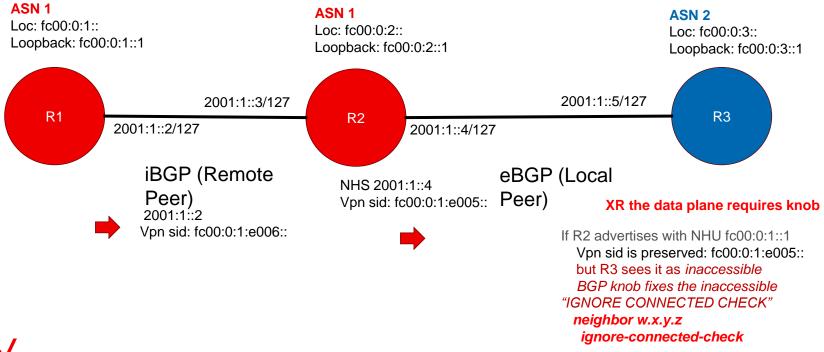
#### **Rule for Inter Domain Peering**

Locator Reachability (That's it!!)

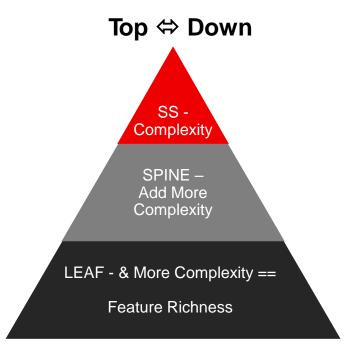
This allows host endpoints to provide static steering capabilities without PCE across any SR Algo cross domain



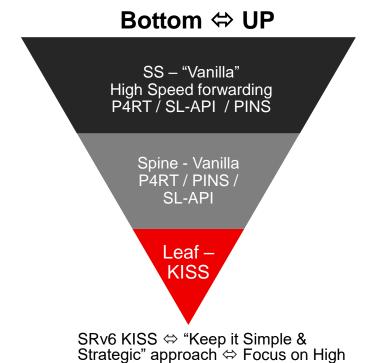
# INTER DOMAIN SRv6 uSID eBGP direct peering (NHU) – (Remote PE)



# SRv6 uSID Design ⇔ "Top ⇔ Down" & "Bottom ⇔ Up" Approach



Traditional mindset has been for feature richness & complexity across the Data Center Fabric



speed forwarding plane packet

pushing throughout

# Real World Use Cases of SRv6 uSID DC Landscape Experience

#### Real World Use Case #1 IPv6 Host Based Networking

Source DC or Domain (IPv6 Only) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain(IPv6 Only) ⇔ Bottom-UP (KISS) [END STATE] or [INTERMEDIATE STATE]

- Traffic Engineering and Carrier Grade features are not a requirement in the Data Center.
- Operators can use white box switches or disaggregated hardware and software with Vanilla IPv6
   Only DC fabric blindly passing the SRv6 uSID packets. Massive bandwidth where Multi Petabits of
   fiber can be thrown at the DC fabric, with the focus on High Bandwidth packet pushing with Ultra
   simplified fabric.
- Steering is initiated from the Data Center host attachment using IGP shortest path leaving the entire fabric 100% vanilla IPv6.



# Real World Use Cases of SRv6 uSID DC Landscape Experience

#### Real World Use Case #2 Dual Plane MPLS / IPv6 Core Migration

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (IPv6 Only) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [INTERMEDIATE STATE]

- Traffic Engineering & Carrier Grade features are a requirement ONLY in the Data Center.
- Traffic Engineering capabilities in the Data Center, and the intermediate domains follow IGP shortest path blindly forwarding the SRv6 uSID packets. Massive scale & resiliency with full carrier grade features in the Data Center.
- Steering is initiated from the DC host attachment and follows IGP shortest path along the intermediate domains to the egress DC or Domain.



# Real World Use Cases of SRv6 uSID DC Landscape Experience

#### Real World Use Case #3 SRv6 uSID End-To-End

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [END STATE]

- Traffic Engineering and all the Carrier Grade features are a requirement.
- Full feature richness.
- **Steering** is initiated from the Data Center host attachment or could be any switch within the DC fabric for any and all flow types.



# SRv6 uSID Host Based Networking Traffic Engineering

### **Options for Host Based Networking**

- eBPF/Cilium (Cilium BGP control plane) CNI or Standalone (Option-1)
- Native Linux Kernel (FRR BGP control plane) –Host Routing with Native Kernel (Option-2)
- Router-in-container (Control Plane & Data Plane) (xRD, SONiC, Nokia, Juniper cRPD) CNF (Option-4)

- Options are listed in order of desirability by operators
- Next few slides we will go into each option in detail



# **Option #1 eBPF/Cilium**

- □ CNI Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNI Provides Traffic Engineering capabilities via SRv6 uSID
- CNI Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

- Data plane programming
- ☐ FRR BGP for bgp control plane advertisement
- ☐ Cilium is one of the most popular CNI's to date and eBPF with its origins in the Linux kernel with its rich policy features & programmability provides seamless integration to compute nodes making it a powerful win-win for developers



# **Option #2 Native Linux Kernel**

- ☐ Host Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ Host Provides Traffic Engineering capabilities via SRv6 uSID
- Host VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

- Data plane programming
- ☐ FRR BGP for control plane advertisement
- ☐ FRR can program the control plane & via Linux Kernel API call program the data plane FIB entries
- ☐ Alternatively, FRR can program the control plane with hook back to Linux Kernel to program the data plane FIB entries



# Option #3 FD.IO VPP

- VPP Connects to global table ⇔ linkage of host fabric to DC fabric
- VPP Provides Traffic Engineering capabilities via SRv6 uSID
- VPP Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

- Data plane programming
- FRR BGP for Control Plane Advertisement
- VPP ceases control of the Linux Hosts NIC
- ☐ Requires Netlink or other method to program VPP FIB
- □ VPP requires application flows to hair-pinned through the Linux Kernel which can become a bottleneck (Not for bandwidth & CPU intensive apps)



# **Option #4 Router-in-Container**

- ☐ CNF Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNF Provides Traffic Engineering capabilities via SRv6 uSID
- CNF Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

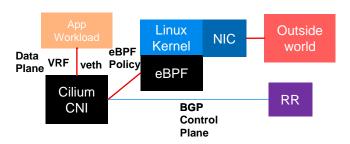
# Router-in-container options (xRD, SONiC, Nokia, Juniper cRPD)

- ☐ Control plane & Data plane programming
- ☐ Requires Linux bridge stitching between Linux Kernel & CNF
- ☐ Container runs in user space so is beneficial for cases where only certain application requires traffic engineering capabilities
- ☐ Very difficult to create linkage of Router in container back to kernel space so all application workloads can take advantage of steering

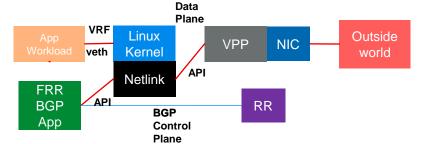


# **Host Networking Drawing**

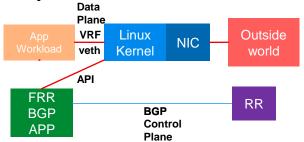
## Option #1 eBPF/Cilium



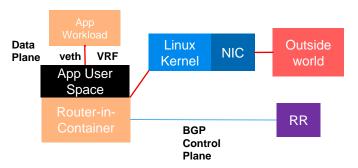
# **Option #3 FD.io VPP**



### **Option #2 Linux Kernel**



# **Option #4 Router-in-Container**





# **Demo time!**

#### All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)
- Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC
  - IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)
- **√**

# Use-Case 1a: SRv6 uSID E2E w/ IPv6 DC & SRv6 uSID

Core Demo time!

All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

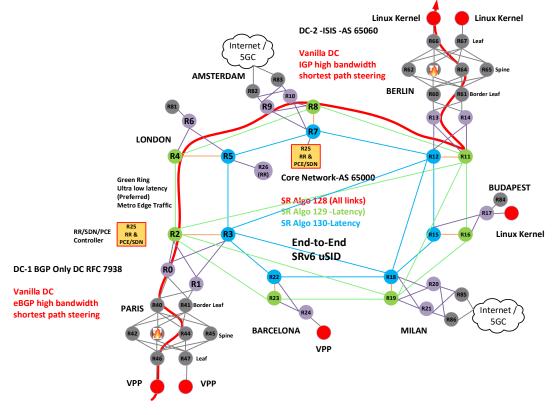
- Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)

- Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC
  - IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)



Use-Case 1a: SRv6 uSID E2E w/ IPv6 DC & SRv6 uSID

Core





# Use-Case 1b: SRv6 uSID E2E w/ SRv6 uSID DC & IPv6

#### Core Demo time!

#### All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

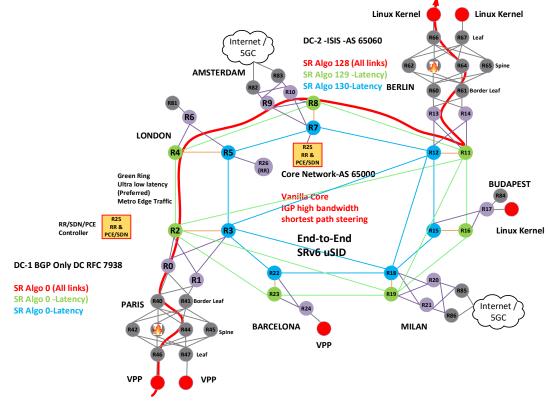
- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- □ Policy programmed from Router-In-Container (XRD)

- Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)



Use-Case 1b: SRv6 uSID E2E w/ SRv6 uSID DC & IPv6

Core





# Use-Case 1c: SRv6 uSID E2E BGP Only & Single AS DC

# **Demo time!**

All Demo's @ YouTube Channel SRv6 uSID DC:

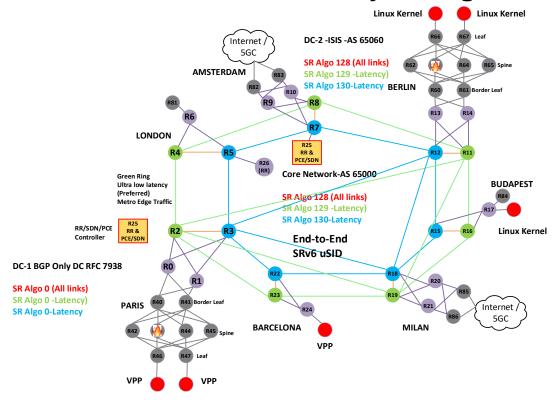
https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- □ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)

- Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC
  - SRv6 uSID End-to-End (1c)



# Use-Case 1c: SRv6 uSID E2E BGP Only & Single AS DC





# **Demo time!**

#### All Demo's @ YouTube Channel IPv6 uSID DC:

https://youtube.com/@SRv6\_uSID\_DC

https://github.com/segmentrouting/srv6-labs

- Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host

#### Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)
- SRv6 uSID uSID End-to-End (1c)

#### Use-case 2: SRv6 uSID w/ Multi POD Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (2a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2b)
- SRv6 uSID End-to-End (2c)

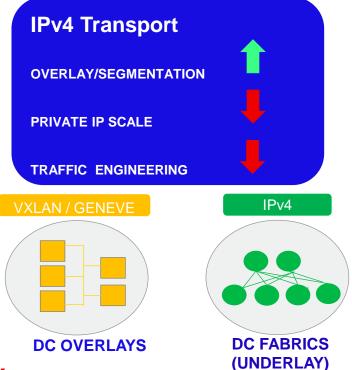
#### Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

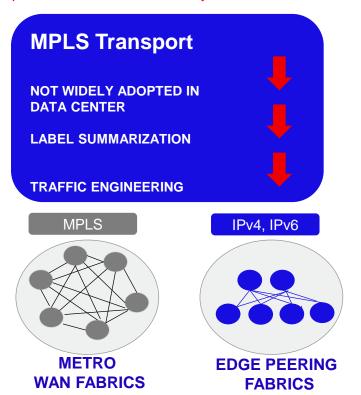
- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)
- SRv6 uSID End-to-End (3c)



# SILOED Networking ⇔ Complexity Tax

Each siloed network Domain has its own Hardware, Software, SDN Stack, Operations & Automation Ecosystem





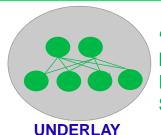
# SRv6 uSID ⇔ Simplicity, Functionality, Ultra Scale

A Common End-to-End Forwarding Architecture Enables Common HW, SW, SDN, Ops ⇔ Massive Scale



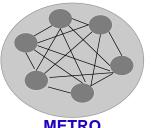
#### SRv6 uSID ⇔ Translates into Ultra Massive Scale



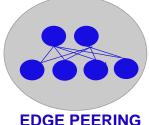


**DC FABRICS** 

"End-to-End Inter-Domain Routing via SRv6 uSID"

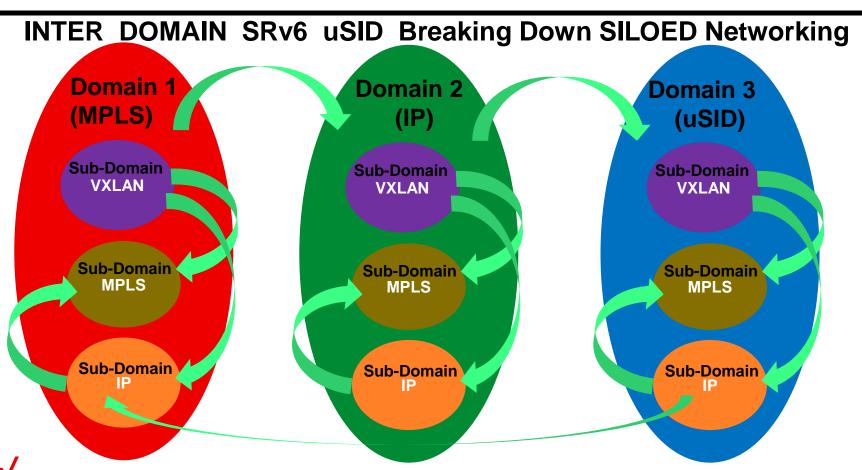


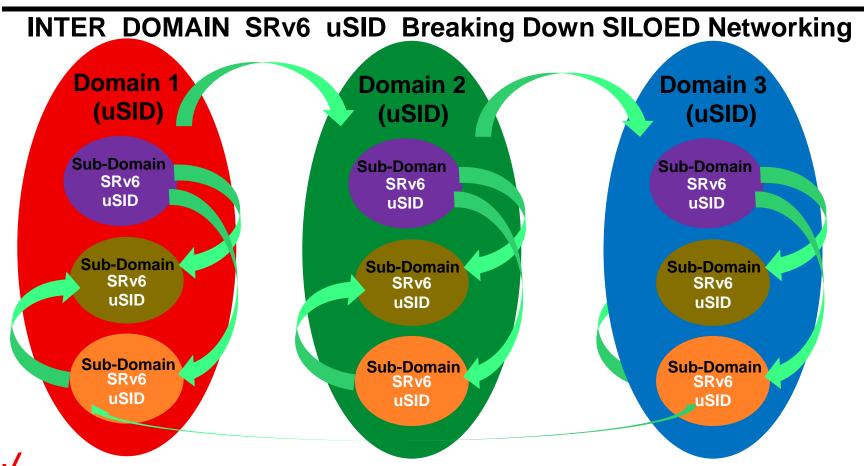
METRO WAN FABRICS



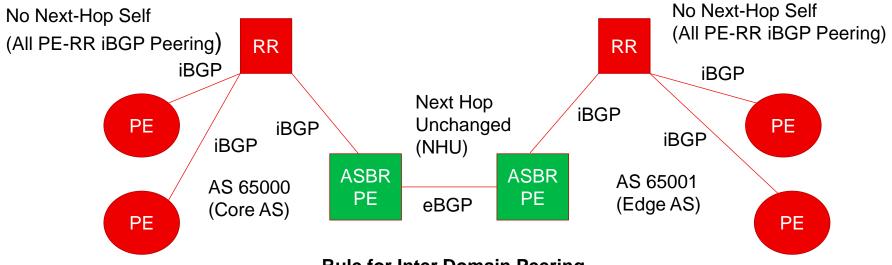
FABRICS







# INTER DOMAIN SRv6 uSID



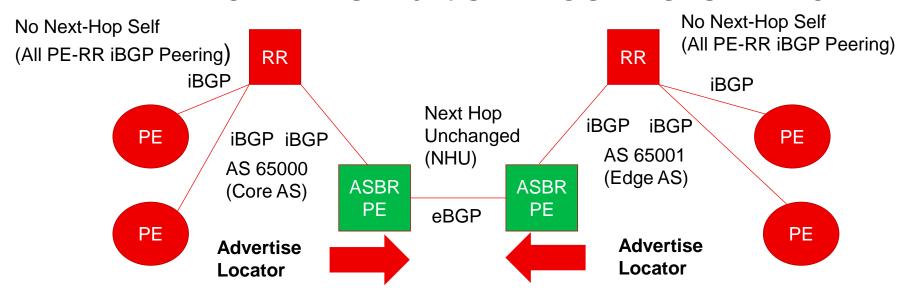
## **Rule for Inter Domain Peering**

- iBGP -No Next-Hop Self
- eBGP –Next Hop Unchanged

(Requirement to Preserve L2 VPN & L3 VPN Service SID across INTER-AS Boundary)



# INTER DOMAIN SRV6 uSID ROUTING SIMPLICITY



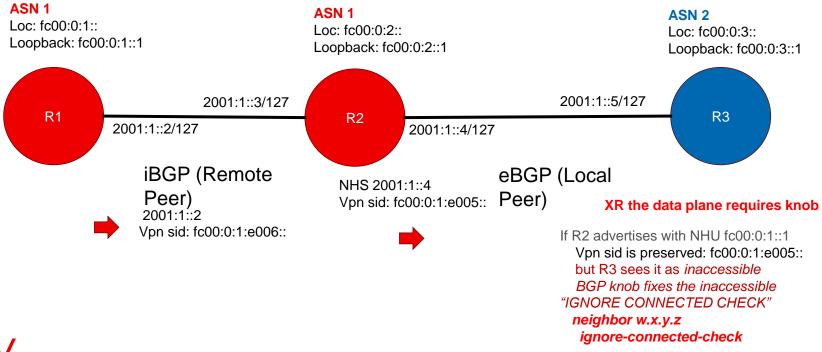
#### **Rule for Inter Domain Peering**

Locator Reachability (That's it!!)

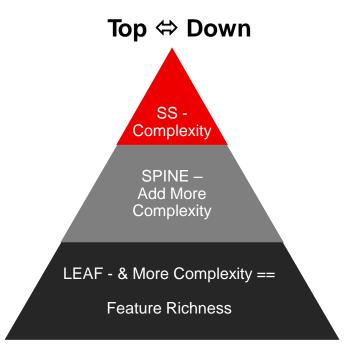
This allows host endpoints to provide static steering capabilities without PCE across any SR Algo cross domain



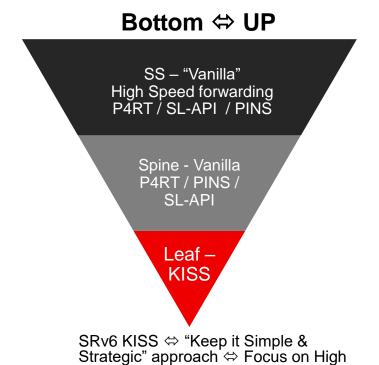
# eBGP direct peering (NHU) – (Remote PE)



# SRv6 uSID Design ⇔ "Top ⇔ Down" & "Bottom ⇔ Up" Approach



Traditional mindset has been for feature richness & complexity across the Data Center Fabric



speed forwarding plane packet

pushing throughout

#### Real World Use Case #1 IPv6 Host Based Networking

Source DC or Domain (IPv6 Only) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain(IPv6 Only) ⇔ Bottom-UP (KISS) [END STATE] or [INTERMEDIATE STATE]

- Traffic Engineering and Carrier Grade features are not a requirement in the Data Center.
- Operators can use white box switches or disaggregated hardware and software with Vanilla IPv6
   Only DC fabric blindly passing the SRv6 uSID packets. Massive bandwidth where Multi Petabits of
   fiber can be thrown at the DC fabric, with the focus on High Bandwidth packet pushing with Ultra
   simplified fabric.
- Steering is initiated from the Data Center host attachment using IGP shortest path leaving the entire fabric 100% vanilla IPv6.



#### Real World Use Case #2 Dual Plane MPLS / IPv6 Core Migration

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (IPv6 Only) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [INTERMEDIATE STATE]

- Traffic Engineering & Carrier Grade features are a requirement ONLY in the Data Center.
- Traffic Engineering capabilities in the Data Center, and the intermediate domains follow IGP shortest path blindly forwarding the SRv6 uSID packets. Massive scale & resiliency with full carrier grade features in the Data Center.
- Steering is initiated from the DC host attachment and follows IGP shortest path along the intermediate domains to the egress DC or Domain.



#### Real World Use Case #3 SRv6 uSID End-To-End

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [END STATE]

- Traffic Engineering and all the Carrier Grade features are a requirement.
- Full feature richness.
- Steering is initiated from the Data Center host attachment or could be any switch within the DC fabric for any and all flow types.



# SRv6 uSID Host Based Networking Traffic Engineering

### **Options for Host Based Networking**

- eBPF/Cilium (Cilium BGP control plane) CNI or Standalone (Option-1)
- Native Linux Kernel (FRR BGP control plane) –Host Routing with Native Kernel (Option-2)
- ☐ FD.io VPP (FRR BGP control plane) -Host Routing via VPP (Option-3)
- Router-in-container (Control Plane & Data Plane) (xRD, SONiC, Nokia, Juniper cRPD) CNF (Option-4)

- Options are listed in order of desirability by operators
- Next few slides we will go into each option in detail



# **Option #1 eBPF/Cilium**

- □ CNI Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNI Provides Traffic Engineering capabilities via SRv6 uSID
- CNI Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for bgp control plane advertisement
- ☐ Cilium is one of the most popular CNI's to date and eBPF with its origins in the Linux kernel with its rich policy features & programmability provides seamless integration to compute nodes making it a powerful win-win for developers



# **Option #2 Native Linux Kernel**

- ☐ Host Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ Host Provides Traffic Engineering capabilities via SRv6 uSID
- Host VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for control plane advertisement
- ☐ FRR can program the control plane & via Linux Kernel API call program the data plane FIB entries
- ☐ Alternatively, FRR can program the control plane with hook back to Linux Kernel to program the data plane FIB entries



# Option #3 FD.IO VPP

- □ VPP Connects to global table ⇔ linkage of host fabric to DC fabric
- VPP Provides Traffic Engineering capabilities via SRv6 uSID
- VPP Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- FRR BGP for Control Plane Advertisement
- VPP ceases control of the Linux Hosts NIC
- ☐ Requires Netlink or other method to program VPP FIB
- □ VPP requires application flows to hair-pinned through the Linux Kernel which can become a bottleneck (Not for bandwidth & CPU intensive apps)



# **Option #4 Router-in-Container**

- ☐ CNF Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNF Provides Traffic Engineering capabilities via SRv6 uSID
- □ CNF Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

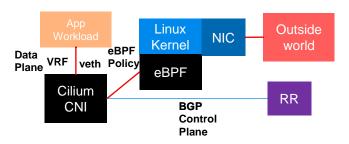
Router-in-container options (xRD, SONiC, Nokia, Juniper cRPD)

- ☐ Control plane & Data plane programming
- ☐ Requires Linux bridge stitching between Linux Kernel & CNF
- ☐ Container runs in user space so is beneficial for cases where only certain application requires traffic engineering capabilities
- ☐ Very difficult to create linkage of Router in container back to kernel space so all application workloads can take advantage of steering

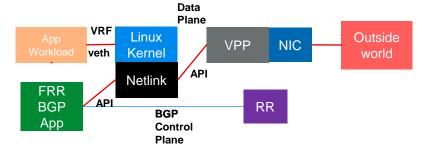


# **Host Networking Drawing**

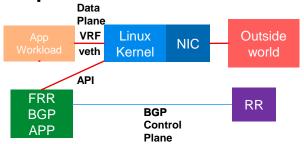
### Option #1 eBPF/Cilium



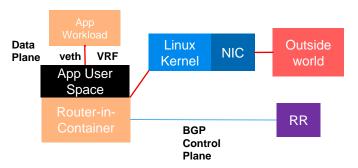
## **Option #3 FD.io VPP**



#### **Option #2 Linux Kernel**



## **Option #4 Router-in-Container**





# **Demo time!**

#### All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)
- Use-case 2: SRv6 uSID with Multi POD Fabrics
  - IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (2a)
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2b)
  - SRv6 uSID End-to-End (2c)



## Use-Case 2a: SRv6 uSID E2E Multi POD w/ IPv6 DC & SRv6 uSID

#### Core Demo time!

All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- □ Policy programmed from Router-In-Container (XRD)

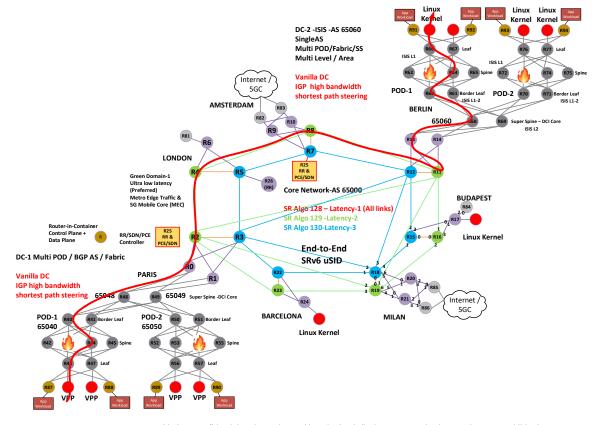
- Use-case 2: SRv6 uSID with Multi POD Fabrics
  - ♦ IPv6 DC 

    ⇒ SRv6 uSID Core 

    ⇒ IPv6 DC (2a)



# Use-Case 2a: SRv6 uSID E2E Multi POD w/ IPv6 DC & SRv6 uSID Core





### Use-Case 2b: SRv6 uSID E2E Multi POD w/ SRv6 uSID DC & IPv6 Core

# **Demo time!**

All Demo's @ YouTube Channel SRv6 uSID DC:

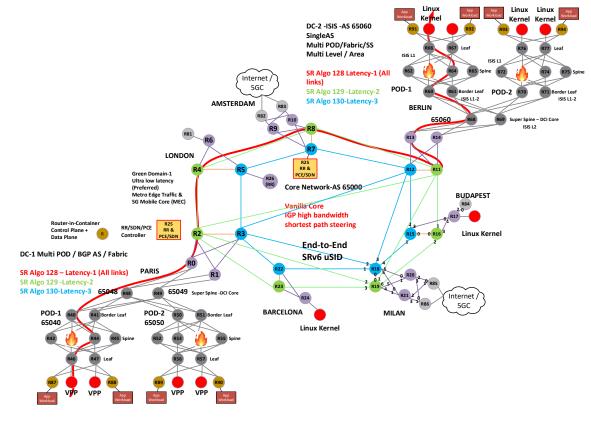
https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- □ Policy programmed from Router-In-Container (XRD)

- Use-case 2: SRv6 uSID with Multi POD Fabrics
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2c)



# Use-Case 2b: SRv6 uSID E2E Multi POD w/ SRv6 uSID DC & IPv6 Core





# Use-Case 2c: SRv6 uSID E2E Multi POD

# **Demo time!**

All Demo's @ YouTube Channel SRv6 uSID DC:

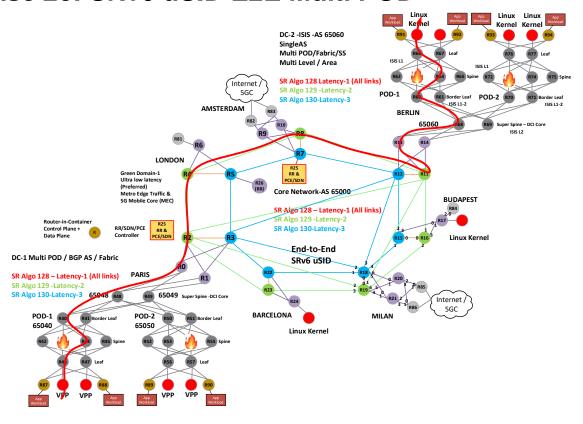
https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)

- Use-case 2: SRv6 uSID with Multi POD Fabrics
  - SRv6 uSID End-to-End (2c)



# Use-Case 2c: SRv6 uSID E2E Multi POD





## **Demo time!**

#### All Demo's @ YouTube Channel IPv6 uSID DC:

https://youtube.com/@SRv6\_uSID\_DC

https://github.com/segmentrouting/srv6-labs

- Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host

#### Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (1a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (1b)
- SRv6 uSID uSID End-to-End (1c)

#### Use-case 2: SRv6 uSID w/ Multi POD Fabrics

- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (2a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (2b)
- SRv6 uSID End-to-End (2c)

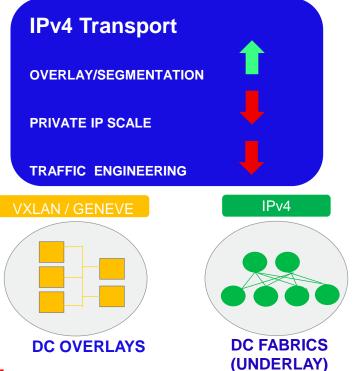
#### Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

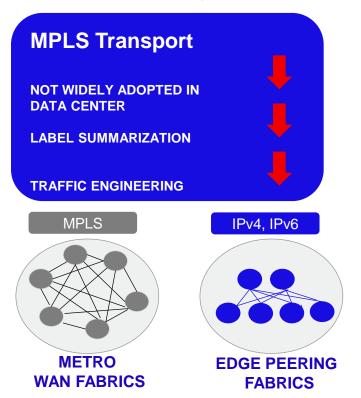
- IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)
- SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)
- SRv6 uSID End-to-End (3c)



# SILOED Networking ⇔ Complexity Tax

Each siloed network Domain has its own Hardware, Software, SDN Stack, Operations & Automation Ecosystem





# SRv6 uSID ⇔ Simplicity, Functionality, Ultra Scale

A Common End-to-End Forwarding Architecture Enables Common HW, SW, SDN, Ops ⇔ Massive Scale



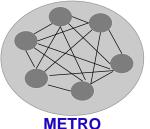
#### SRv6 uSID ⇔ Translates into Ultra Massive Scale



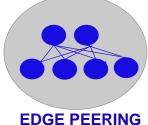


**DC FABRICS** 

"End-to-End Inter-Domain Routing via SRv6 uSID"

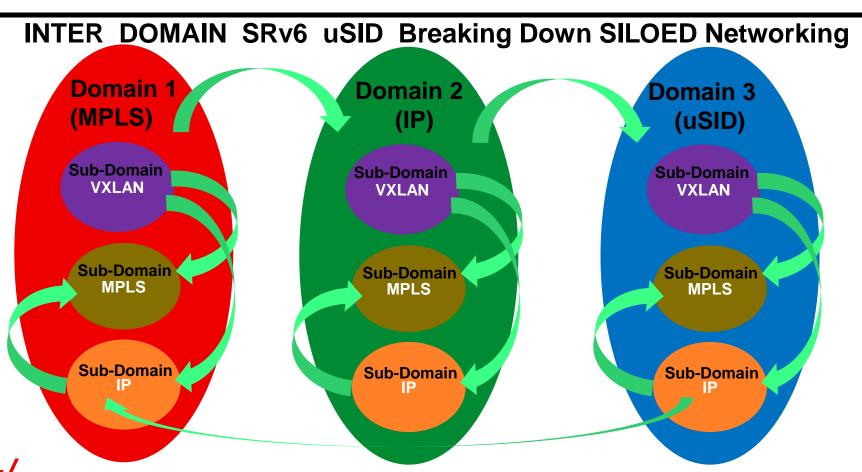


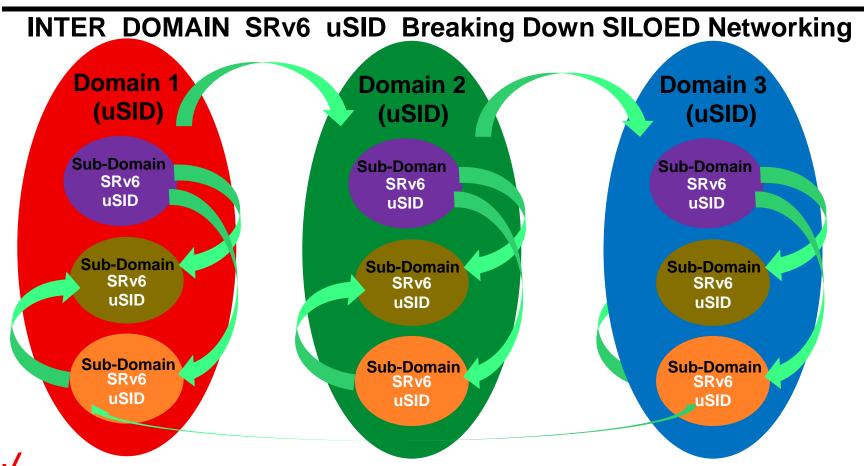
METRO WAN FABRICS



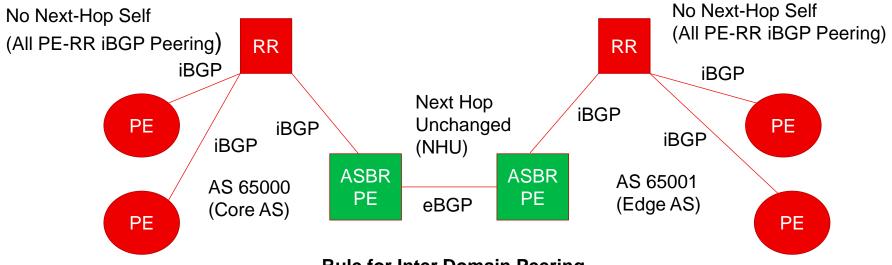
FABRICS







# INTER DOMAIN SRv6 uSID



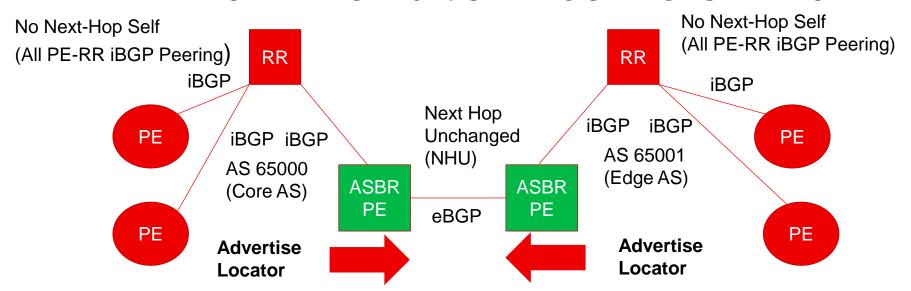
## Rule for Inter Domain Peering

- iBGP -No Next-Hop Self
- eBGP –Next Hop Unchanged

(Requirement to Preserve L2 VPN & L3 VPN Service SID across INTER-AS Boundary)



# INTER DOMAIN SRV6 uSID ROUTING SIMPLICITY



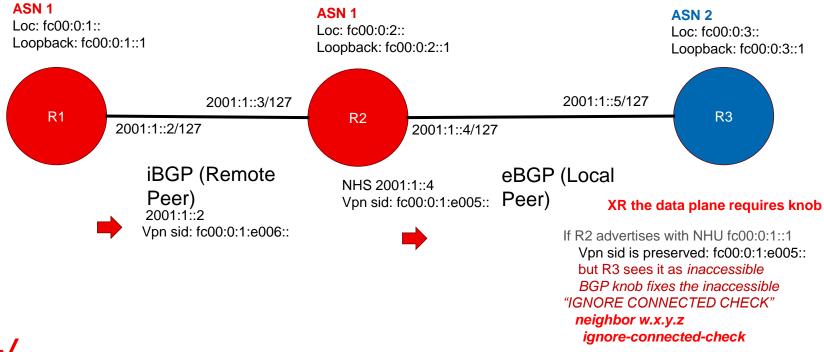
#### **Rule for Inter Domain Peering**

Locator Reachability (That's it!!)

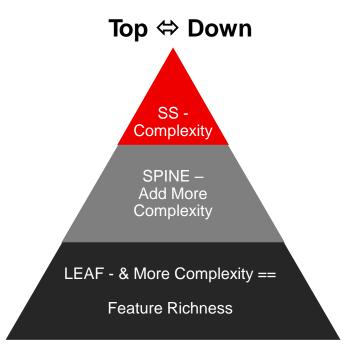
This allows host endpoints to provide static steering capabilities without PCE across any SR Algo cross domain



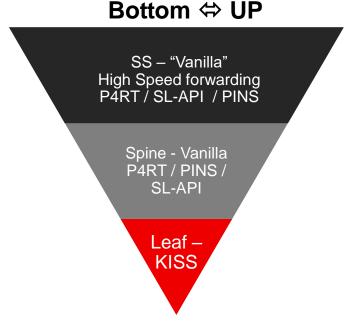
# eBGP direct peering (NHU) – (Remote PE)



# SRv6 uSID Design ⇔ "Top ⇔ Down" & "Bottom ⇔ Up" Approach



Traditional mindset has been for feature richness & complexity across the Data Center Fabric



SRv6 KISS ⇔ "Keep it Simple & Strategic" approach ⇔ Focus on High speed forwarding plane packet pushing throughput



#### Real World Use Case #1 IPv6 Host Based Networking

Source DC or Domain (IPv6 Only) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain(IPv6 Only) ⇔ Bottom-UP (KISS) [END STATE] or [INTERMEDIATE STATE]

- Traffic Engineering and Carrier Grade features are not a requirement in the Data Center.
- Operators can use white box switches or disaggregated hardware and software with Vanilla IPv6
   Only DC fabric blindly passing the SRv6 uSID packets. Massive bandwidth where Multi Petabits of
   fiber can be thrown at the DC fabric, with the focus on High Bandwidth packet pushing with Ultra
   simplified fabric.
- Steering is initiated from the Data Center host attachment using IGP shortest path leaving the entire fabric 100% vanilla IPv6.



#### Real World Use Case #2 Dual Plane MPLS / IPv6 Core Migration

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (IPv6 Only) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [INTERMEDIATE STATE]

- Traffic Engineering & Carrier Grade features are a requirement ONLY in the Data Center.
- Traffic Engineering capabilities in the Data Center, and the intermediate domains follow IGP shortest path blindly forwarding the SRv6 uSID packets. Massive scale & resiliency with full carrier grade features in the Data Center.
- Steering is initiated from the DC host attachment and follows IGP shortest path along the intermediate domains to the egress DC or Domain.



#### Real World Use Case #3 SRv6 uSID End-To-End

Source DC or Domain (SRv6 uSID Capable) ⇔ Core Domain / Multi Domain OAD (One Administrative Domain) / Internet (SRv6 uSID Capable) ⇔ Destination Data Center or Domain (SRv6 uSID Capable) ⇔ Top-Down Feature Rich [END STATE]

- Traffic Engineering and all the Carrier Grade features are a requirement.
- Full feature richness.
- Steering is initiated from the Data Center host attachment or could be any switch within the DC fabric for any and all flow types.



# SRv6 uSID Host Based Networking Traffic Engineering

### **Options for Host Based Networking**

- eBPF/Cilium (Cilium BGP control plane) CNI or Standalone (Option-1)
- Native Linux Kernel (FRR BGP control plane) –Host Routing with Native Kernel (Option-2)
- ☐ FD.io VPP (FRR BGP control plane) -Host Routing via VPP (Option-3)
- ☐ Router-in-container (Control Plane & Data Plane) (xRD, SONiC, Nokia, Juniper cRPD) CNF (Option-4)

- Options are listed in order of desirability by operators
- Next few slides we will go into each option in detail



# **Option #1 eBPF/Cilium**

- □ CNI Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNI Provides Traffic Engineering capabilities via SRv6 uSID
- CNI Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for bgp control plane advertisement
- ☐ Cilium is one of the most popular CNI's to date and eBPF with its origins in the Linux kernel with its rich policy features & programmability provides seamless integration to compute nodes making it a powerful win-win for developers



# **Option #2 Native Linux Kernel**

- ☐ Host Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ Host Provides Traffic Engineering capabilities via SRv6 uSID
- Host VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for control plane advertisement
- ☐ FRR can program the control plane & via Linux Kernel API call program the data plane FIB entries
- ☐ Alternatively, FRR can program the control plane with hook back to Linux Kernel to program the data plane FIB entries



## Option #3 FD.IO VPP

- □ VPP Connects to global table ⇔ linkage of host fabric to DC fabric
- VPP Provides Traffic Engineering capabilities via SRv6 uSID
- VPP Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- FRR BGP for Control Plane Advertisement
- VPP ceases control of the Linux Hosts NIC
- Requires Netlink or other method to program VPP FIB
- □ VPP requires application flows to hair-pinned through the Linux Kernel which can become a bottleneck (Not for bandwidth & CPU intensive apps)



# **Option #4 Router-in-Container**

- ☐ CNF Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNF Provides Traffic Engineering capabilities via SRv6 uSID
- CNF Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

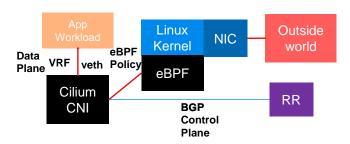
#### Router-in-container options (xRD, SONiC, Nokia, Juniper cRPD)

- ☐ Control plane & Data plane programming
- ☐ Requires Linux bridge stitching between Linux Kernel & CNF
- ☐ Container runs in user space so is beneficial for cases where only certain application requires traffic engineering capabilities
- ☐ Very difficult to create linkage of Router in container back to kernel space so all application workloads can take advantage of steering

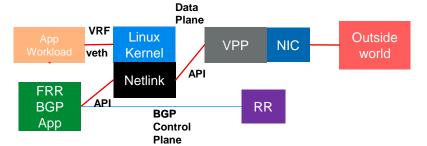


# **Host Networking Drawing**

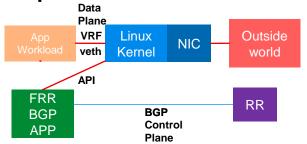
#### Option #1 eBPF/Cilium



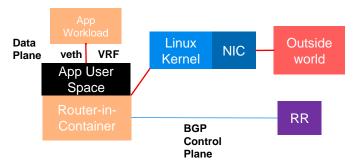
#### **Option #3 FD.io VPP**



#### **Option #2 Linux Kernel**



#### **Option #4 Router-in-Container**





#### **Demo time!**

All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs

https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)
- Use-case 3: SRv6 uSID with Multi POD? Domain Fabrics
  - IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)
  - SRv6 uSID End-to-End (3c)



#### Use-Case 3a: SRv6 uSID E2E Multi POD/Domain w/ IPv6 DC & SRv6 uSID

Core Demo time!

All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

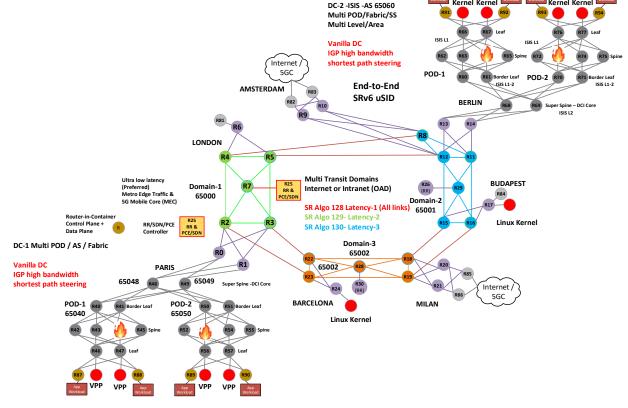
- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)
- Use-case 3: SRv6 uSID with Multi POD? Domain Fabrics
  - IPv6 DC ⇔ SRv6 uSID Core ⇔IPv6 DC (3a)



Use-Case 3a: SRv6 uSID E2E Multi POD/Domain w/ IPv6 DC & SRv6 uSID

Kernel Kernel

Core





### Use-Case 3b: SRv6 uSID E2E Multi POD/Domain w/ SRv6 uSID DC & IPv6 Core

**Demo time!** 

All Demo's @ YouTube Channel SRv6 uSID DC:

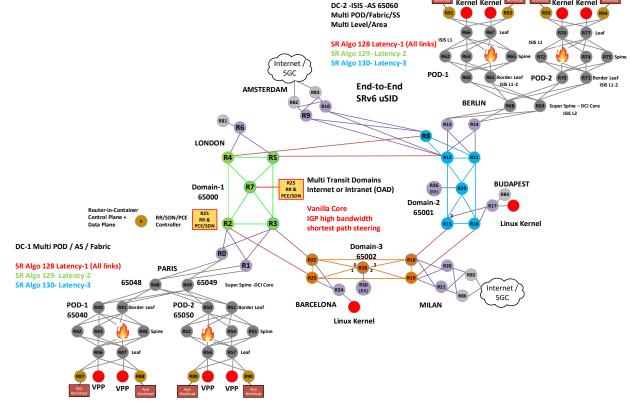
https://github.com/segmentrouting/srv6-labs

https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- Policy programmed from Router-In-Container (XRD)
- Use-case 3: SRv6 uSID with Multi POD? Domain Fabrics
  - SRv6 uSID DC ⇔ IPv6 Core ⇔ SRv6 uSID DC (3b)



Use-Case 3b: SRv6 uSID E2E Multi POD/Domain w/ SRv6 uSID DC & IPv6 Core





#### Use-Case 3c: SRv6 uSID E2E Multi POD / Multi Domain

#### **Demo time!**

All Demo's @ YouTube Channel SRv6 uSID DC:

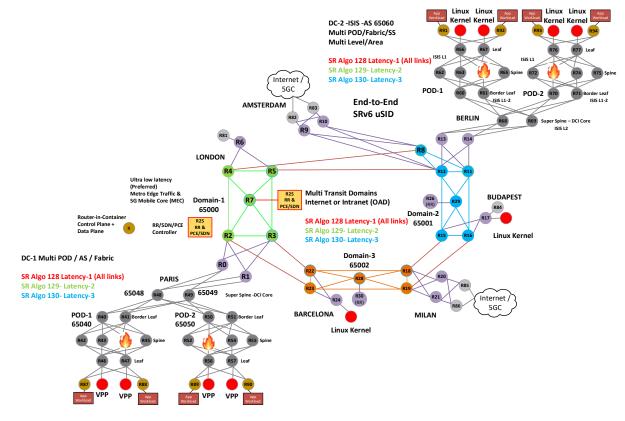
https://github.com/segmentrouting/srv6-labs https://youtube.com/@SRv6\_uSID\_DC

- ☐ Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host
- → Policy programmed from Router-In-Container (XRD)

- Use-case 3: SRv6 uSID with Multi POD? Domain Fabrics
  - SRv6 uSID End-to-End (3c)



#### Use-Case 3c: SRv6 uSID E2E Multi POD / Multi Domain





# Use-Case 1: SRv6 uSID DC Fabric Packet Capture & Screen

SCIA Sessib IPv4 payload steer DC-2 Paris  $\Leftrightarrow$  Core  $\Leftrightarrow$  DC-1 Berlin using VPP Host attached to DC fabric

#### SRv6 uSID IPv4 payload steering policy

vpp#sr policy add bsid 40::40 next fc00:0:44:40:4:64:66:e000 encap

vpp#sr steer I3 10.0.0.66/32 via bsid 40::40

vpp# show sr policies SR policies:

[0].- BSID: 40::40

Behavior: Encapsulation EncapSrcIP: fc00:0:46:1::3

Type: Default FIB table: 0 Segment Lists:

[0].- < fc00:0:44:40:4:64:66:e000 > weight: 1

vpp# show sr steering-policies SR steering policies: Traffic SR policy BSID L3 10.0 0.66/32 40::40

#### SRv6 uSID IPv6 payload steer:

vpp#sr policy add bsid 41::41 next fc00:0:44:40:4:64:66:e000 encap

vpp# sr steer l3 fc00:0:66::1/128 via bsid 41::41

vpp# show sr policies

SR policies:

[1].- BSID: 94::94

Behavior: Encapsulation EncapSrcIP: fc00:0:46:1::3

Type: Default FIB table: 0 Segment Lists:

[1].- < fc00:0:45:41:66:e000:: > weight: 1

.....

vpp# show sr steering-policies

SR steering policies:

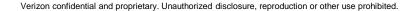
Traffic SR policy BSID L3 fc00:0:66::1/128 41::41

#### IPv4 payload packet capture xrd61-xrd64

04:41:56.724814 IP6 fc00:0:46:1::3 > fc00:0:64:66:e000::: IP 10.11.46.2 > 10.0.0.66: ICMP echo request, id 113, seq 463, length 64 04:41:57.724271 IP6 fc00:0:46:1::3 > fc00:0:64:66:e000::: IP 10.11.46.2 > 10.0.0.66: ICMP echo request, id 113, seq 464, length 64

#### IPv6 payload packet capture xrd61-xrd64:

04:55:37.285381 IP6 fc00:0:46:1::3 > fc00:0:64:66:e000::: IP6 fc00:0:46:2::2 > fc00:0:66::1: ICMP6, echo request, seq 368, length 64 04:55:38.285012 IP6 fc00:0:46:1::3 > fc00:0:64:66:e000::: IP6 fc00:0:46:2::2 > fc00:0:66::1: ICMP6, echo request, seq 369, length 64



# Use-Case 2: SRv6 uSID Inter-DC Packet Capture & Screen

Scrapes
SRV6 usin PV4 payload steer DC-1 Berlin  $\Leftrightarrow$  Core  $\Leftrightarrow$  DC-2 Paris using Linux host attached to DC fabric

#### SRv6 uSID IPv6 payload steering policy:

root@ubuntu-linux-srv6:sudo ip route add 10.0.0.46/32 encap seg6 mode encap segs fc00:0:64:61:4:44:46:e000 dev ens7

root@ubuntu-linux-srv6:/home/cisco# ip route

default via 192.168.122.1 dev ens8 proto dhcp src 192.168.122.88 metric 100

10.0.0.0/24 via 10.10.66.2 dev ens7 proto static

10.0.0.46 encap seg6 mode encap segs 1 [fc00:0:64:61:4:44:46:e000 ] dev ens7 scope link------>SRv6 uSID steering programmed IPv4 payload

#### SRv6 uSID IPv4 payload steer capture DC-2 Paris:

xrd41-xrd44

#### SRv6 uSID IPv6 payload steering policy:

root@ubuntu-linux-srv6:sudo ip -6 route add fc00:0:46::1 encap seg6 mode encap segs fc00:0:64:61:4:44:46:e000 dev ens7

root@ubuntu-linux-srv6:/home/cisco# ip -6 route

::1 dev lo proto kernel metric 256 pref medium

fc00:0:46::1 encap seg6 mode encap segs 1 [fc00:0:64:61:4:44:46:e000 ] dev ens7 metric 1024 pref medium--->SRv6 uSID steering programmed IPv6 payload

#### SRv6 uSID IPv6 payload steer capture DC-2 Paris:

xrd44-xrd46

20:40:32.890109 IP6 fc00:0:66:1:5054:2ff:fe41:b107 > fc00:0:46:e000::: srcrt (len=2, type=4, segleft=0[|srcrt] 20:40:32.890994 IP6 fc00:0:46::1 > fc00:0:66:1:5054:2ff:fe41:b107: ICMP6, parameter problem, code-#4, length 176



# Design Experience of an SRv6 uSID Data Center

Gyan Mishra

IT Technologist & Innovation Specialist

Associate-Fellow – Network Design

April 9<sup>th</sup> 2024



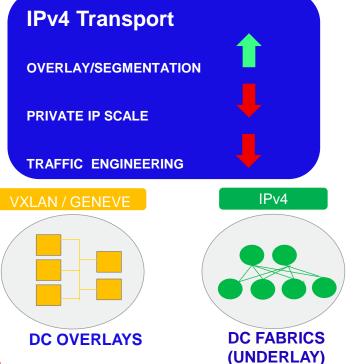
# Objectives for Todays SRv6 uSID Data Center Presentation

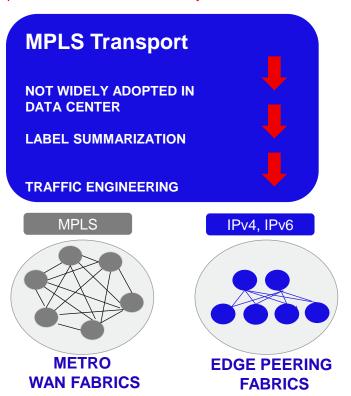
- Building on the foundation from last year
- Intricate design experience details of an SRv6 uSID Data Center
- Review challenges of the innovative solution
- Highlight the transformational potential for the technology
- Comparative analysis to spotlight the efficiencies and performance enhancements outpacing todays technologies at record speed
- Navigate real world design experiences and possible tangible outcomes to comprehend the sheer power of SRv6 uSID
- In the end the goal is to underscore the immense value proposition with SRv6 uSID in the Data Center landscape



# SILOED Networking ⇔ Complexity Tax

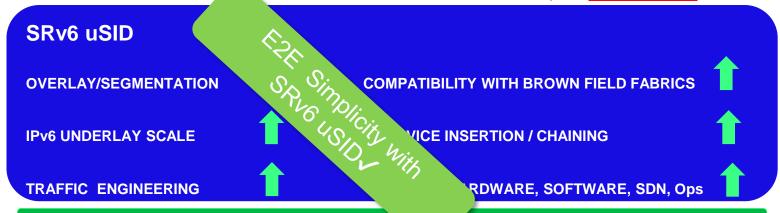
Each siloed network Domain has its own Hardware, Software, SDN Stack, Operations & Automation Ecosystem



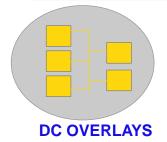


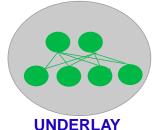
# SRv6 uSID ⇔ Simplicity, Functionality, Ultra Scale

A Common End-to-End Forwardin hitecture Enables Common HW, SW, SDN, Ops ⇔ Massive Scale



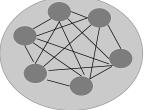
SRv6 uSID ⇔ Translates . A Massive Scale





**DC FABRICS** 

"End-to-End Inter-Domain Routing via SRv6 uSID"



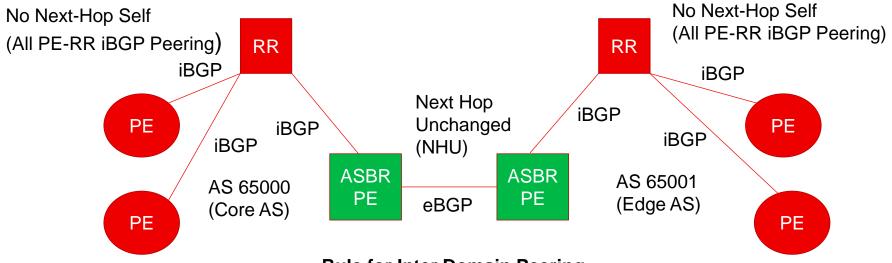
METRO WAN FABRICS



EDGE PEERING FABRICS



# INTER DOMAIN SRv6 uSID



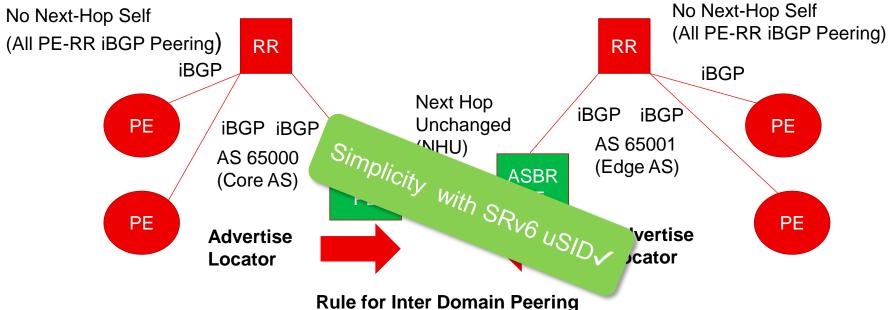
#### Rule for Inter Domain Peering

- iBGP -No Next-Hop Self
- eBGP –Next Hop Unchanged

(Requirement to Preserve L2 VPN & L3 VPN Service SID across INTER-AS Boundary)



# INTER DOMAIN SRV6 USID ROUTING SIMPLICITY



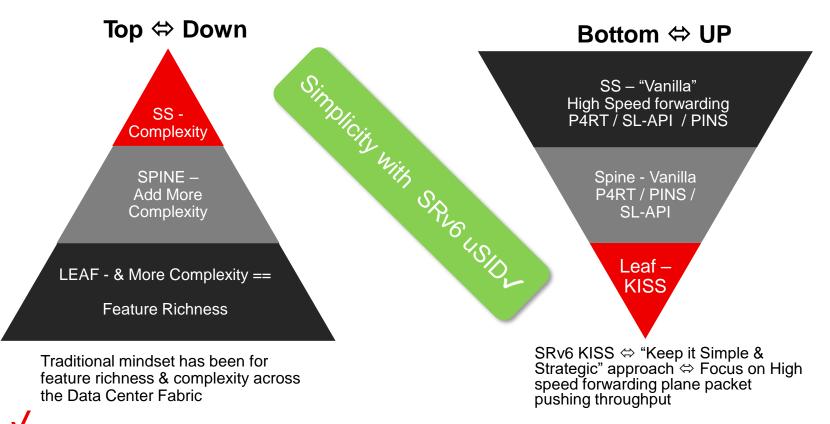
#### Rule for Inter Domain Peering

Locator Reachability (That's it!!)

This allows host endpoints to provide static steering capabilities without PCE across any SR Algo cross domain



# SRv6 uSID Design ⇔ "Top ⇔ Down" & "Bottom ⇔ Up" Approach



# Real World Use-cases

- **#1 IPv6 Host Based Networking**
- #2 Dual Plane MPLS / IPv6 Core Migration
- #3 SRv6 uSID End-To-End



# **#1 IPv6 Host Based Networking**

- Traffic Engineering and Carrier Grade features are not a requirement in the Data Center.
- Operators can use white box switches or disaggregated hardware and software with Vanilla IPv6 Only DC fabric blindly passing the SRv6 uSID packets. Massive bandwidth where Multi Petabits of fiber can be thrown at the DC fabric, with the focus on High Bandwidth packet pushing with Ultra simplified fabric.
- **Steering** is initiated from the Data Center host attachment using IGP shortest path leaving the entire fabric 100% vanilla IPv6.



# #2 Dual Plane MPLS / IPv6 Core Migration

- Traffic Engineering & Carrier Grade features are a requirement ONLY in the Data Center.
- Traffic Engineering capabilities in the Data Center, and the intermediate domains follow IGP shortest path blindly forwarding the SRv6 uSID packets. Massive scale & resiliency with full carrier grade features in the Data Center.
- **Steering** is initiated from the DC host attachment and follows IGP shortest path along the intermediate domains to the egress DC or Domain.

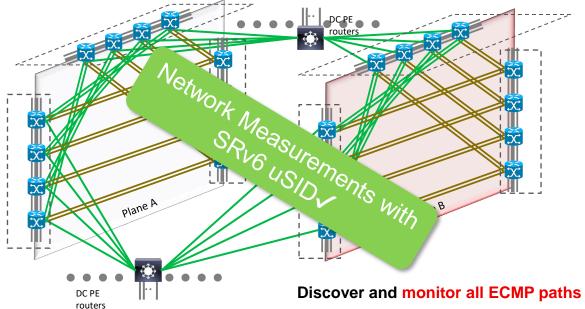


# #3 SRv6 uSID End-To-End

- Traffic Engineering and all the Carrier Grade features are a requirement.
- Full feature richness.
- **Steering** is initiated from the Data Center host attachment or could be any switch within the DC fabric for any and all flow types.



# **Integrated Performance Monitoring (IPM)**

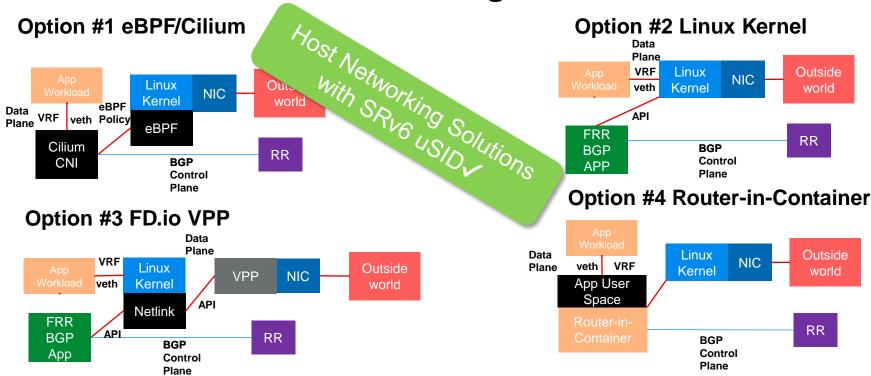


Provide **Enough PPS** to measure all ECMP paths

Report accurately across paths



# **Host Networking Stacks**





# **Simplified Host Networking**

#### **Lightweight Host Routing:**

Linux include its SRv6 Locator (IF fix) within the LLDP advertisements

TOR (IOS XR/SONiC) redistrib

#### **Simpler solution:**

Provides reachability (routing) up to the host

Provides visibility into the container (workload IP address)

Provides liveness detection (built-into LLDP)

→ SRv6 Locator: fc00:0:11::/48

fc00:0:11::/48 via TOR1
fc00:0:12::/48 via TOR1

LLDP – SRv6
Locator:
fc00:0:12::/48

H11 H12

ISIS / BGP:

**TOR** 

Lightweight: No need to run BGP stack on the host



#### **Demo time!**

#### All Demo's @ YouTube Channel SRv6 uSID DC:

https://github.com/segmentrouting/srv6-labs

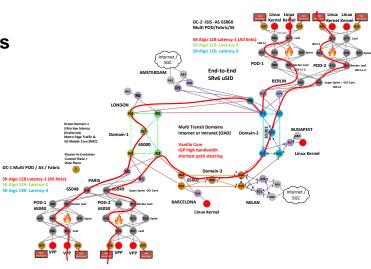
https://youtube.com/@SRv6\_uSID\_DC

Use-case 1: SRv6 uSID BGP-Only DC & Single AS DC

Use-case 2: SRv6 uSID w/ Multi POD Fabrics

Use-case 3: SRv6 uSID w/ Multi POD/Domain Fabrics

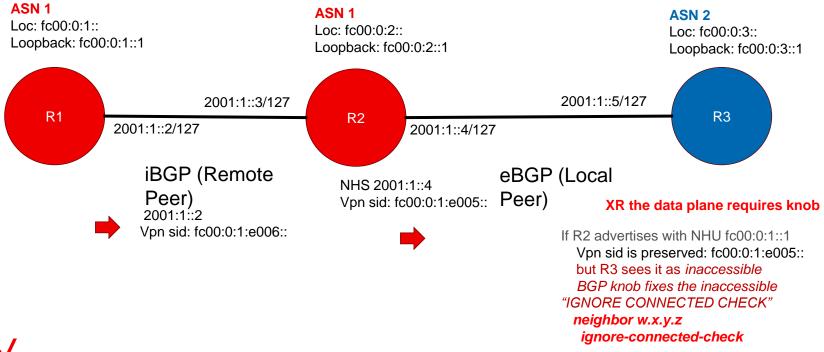
- Host-to-Host multi-pod across the metro
- ☐ Policy programmed from Linux Kernel & VPP host

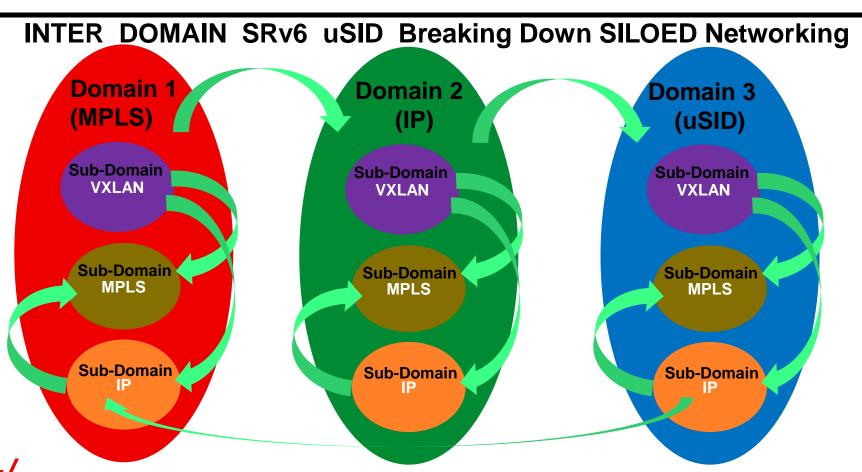


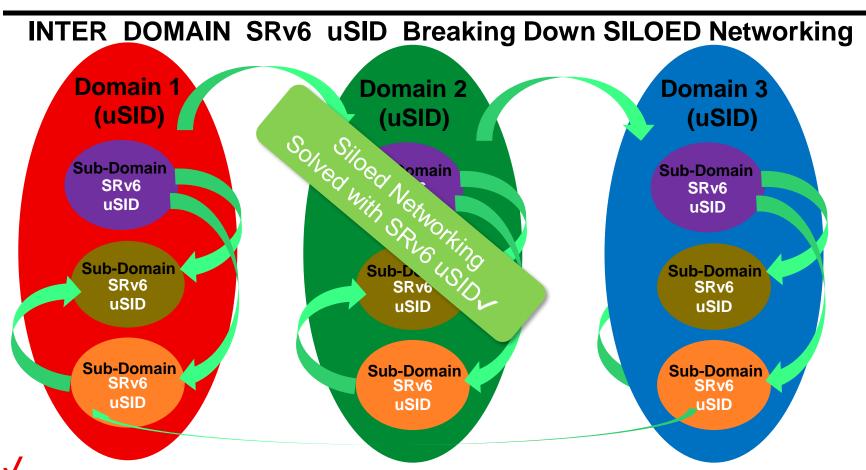


# verizon /

# INTER DOMAIN SRv6 uSID eBGP direct peering (NHU) – (Remote PE)







# SRv6 uSID Host Based Networking Traffic Engineering

#### **Options for Host Based Networking**

- eBPF/Cilium (Cilium BGP control plane) CNI or Standalone (Option-1)
- Native Linux Kernel (FRR BGP control plane) –Host Routing with Native Kernel (Option-2)
- Router-in-container (Control Plane & Data Plane) (xRD, SONiC, Nokia, Juniper cRPD) CNF (Option-4)

- Options are listed in order of desirability by operators
- Next few slides we will go into each option in detail



# **Option #1 eBPF/Cilium**

- □ CNI Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNI Provides Traffic Engineering capabilities via SRv6 uSID
- CNI Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for bgp control plane advertisement
- ☐ Cilium is one of the most popular CNI's to date and eBPF with its origins in the Linux kernel with its rich policy features & programmability provides seamless integration to compute nodes making it a powerful win-win for developers



# **Option #2 Native Linux Kernel**

- ☐ Host Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ Host Provides Traffic Engineering capabilities via SRv6 uSID
- Host VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- ☐ FRR BGP for control plane advertisement
- ☐ FRR can program the control plane & via Linux Kernel API call program the data plane FIB entries
- ☐ Alternatively, FRR can program the control plane with hook back to Linux Kernel to program the data plane FIB entries



## Option #3 FD.IO VPP

- □ VPP Connects to global table ⇔ linkage of host fabric to DC fabric
- VPP Provides Traffic Engineering capabilities via SRv6 uSID
- VPP Provides VPN overlay Workload Container, VM, CNF, VNF

- Data plane programming
- FRR BGP for Control Plane Advertisement
- VPP ceases control of the Linux Hosts NIC
- Requires Netlink or other method to program VPP FIB
- □ VPP requires application flows to hair-pinned through the Linux Kernel which can become a bottleneck (Not for bandwidth & CPU intensive apps)



# **Option #4 Router-in-Container**

- ☐ CNF Connects to global table ⇔ linkage of host fabric to DC fabric
- ☐ CNF Provides Traffic Engineering capabilities via SRv6 uSID
- CNF Provides VPN overlay Workload Container, VM, CNF, VNF

#### **Details**

Router-in-container options (xRD, SONiC, Nokia, Juniper cRPD)

- ☐ Control plane & Data plane programming
- ☐ Requires Linux bridge stitching between Linux Kernel & CNF
- ☐ Container runs in user space so is beneficial for cases where only certain application requires traffic engineering capabilities
- ☐ Very difficult to create linkage of Router in container back to kernel space so all application workloads can take advantage of steering

